

lab3

In [73]:

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_score, cross_validate
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score, classification_report
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.metrics import ConfusionMatrixDisplay

import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(style="ticks")
```

In [74]:

```
### 1. Выберите набор данных (датасет) для решения задачи
классификации или регрессии.
```

In [75]:

```
data = pd.read_csv('Automobile.csv', sep=",")
```

In [76]:

```
data.shape
```

Out[76]:

```
(398, 9)
```

In [77]:

```
data.dtypes
```

Out[77]:

name	object
mpg	float64
cylinders	int64
displacement	float64
horsepower	int64
weight	int64
acceleration	float64

```
model_year      int64
origin          object
dtype: object
```

In [78]:

```
data.isnull().sum()
```

Out[78]:

```
name          0
mpg           0
cylinders     0
displacement  0
horsepower    0
weight        0
acceleration  0
model_year    0
origin        0
dtype: int64
```

In [79]:

2. В случае необходимости проведите удаление или заполнение пропусков и кодирование категориальных признаков.

In [80]:

```
data.head()
```

Out[80]:

	name	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin
0	chevrolet chevelle malibu	18.0	8	307.0	130	3504	12.0	70	usa
1	buick skylark 320	15.0	8	350.0	165	3693	11.5	70	usa
2	plymouth satellite	18.0	8	318.0	150	3436	11.0	70	usa

	name	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin
3	amc rebel sst	16.0	8	304. 0	150	3433	12.0	70	usa
4	ford torin o	17.0	8	302. 0	140	3449	10.5	70	usa

In [81]:

```
data["cylinders"].unique()
```

Out[81]:

```
array([8, 4, 6, 3, 5], dtype=int64)
```

In [82]:

```
# Удаление колонок
```

```
data.drop(columns=['name'], inplace=True)
```

```
data.drop(columns=['origin'], inplace=True)
```

```
# Перемещение колонки "cylinders" в начало датафрейма
```

```
seed_type_column = data['cylinders']
```

```
data.drop(columns=['cylinders'], inplace=True)
```

```
data.insert(0, 'cylinders', seed_type_column)
```

```
data.head()
```

Out[82]:

	cylinders	mpg	displacement	horsepower	weight	acceleration	model_year
0	8	18.0	307.0	130	3504	12.0	70
1	8	15.0	350.0	165	3693	11.5	70
2	8	18.0	318.0	150	3436	11.0	70
3	8	16.0	304.0	150	3433	12.0	70
4	8	17.0	302.0	140	3449	10.5	70

In [83]:

```
fig, ax = plt.subplots(1, 1, sharex='col', sharey='row',
figsize=(20,10))
```

```
fig.suptitle('Корреляционная матрица')
```

```
sns.heatmap(data.corr(), ax=ax, annot=True, fmt='.3f')
```

Out[83]:

<Axes: >

No description has been provided for this image

In [84]:

```
fig, ax = plt.subplots()

color_dict = {4: 'r', 6: 'g', 8: 'b'}

for cylinders, color in color_dict.items():
    data_class = data[data['cylinders'] == cylinders]

    ax.scatter(data_class['acceleration'], data_class['horsepower'],
               color=color,
               label=f'Cylinders: {cylinders}')

ax.set_xlabel('acceleration')
ax.set_ylabel('horsepower')
ax.legend()

plt.show()
```

No description has been provided for this image

In [85]:

```
data_x = data.copy()
data_x = data_x.drop(columns="cylinders")
data_x
```

Out[85]:

	mpg	displace ment	horsepo wer	weight	accelera tion	model_y ear
0	18.0	307.0	130	3504	12.0	70
1	15.0	350.0	165	3693	11.5	70
2	18.0	318.0	150	3436	11.0	70
3	16.0	304.0	150	3433	12.0	70
4	17.0	302.0	140	3449	10.5	70
...
393	27.0	140.0	86	2790	15.6	82
394	44.0	97.0	52	2130	24.6	82
395	32.0	135.0	84	2295	11.6	82
396	28.0	120.0	79	2625	18.6	82

	mpg	displacement	horsepower	weight	acceleration	model_year
397	31.0	119.0	82	2720	19.4	82

398 rows × 6 columns

In [86]:

```
data_y = data["cylinders"]
data_y
```

Out[86]:

```
0      8
1      8
2      8
3      8
4      8
..
393    4
394    4
395    4
396    4
397    4
Name: cylinders, Length: 398, dtype: int64
```

In [87]:

```
# Разделение выборки на обучающую и тестовую
data_X_test, data_X_train, data_Y_test, data_Y_train =
train_test_split(data_x, data_y, test_size=0.2, random_state=2024)
```

In [88]:

```
# Обучение модели для произвольного гиперпараметра K=50
```

```
cl1_1 = KNeighborsClassifier(n_neighbors=50)
```

```
cl1_1.fit(data_X_train, data_Y_train)
target1_0 = cl1_1.predict(data_X_train)
target1_1 = cl1_1.predict(data_X_test)
```

In [89]:

```
# Оценка качества
accuracy_score(data_Y_train, target1_0), accuracy_score(data_Y_test,
target1_1)
```

Out[89]:

```
(0.7625, 0.7610062893081762)
```

In [90]:

```
scoring = {'precision': 'precision_weighted',
           'recall': 'recall_weighted',
           'f1': 'f1_weighted'}
```

In [91]:

```
confusion_matrix(data_Y_test, target1_1)
```

Out[91]:

```
array([[ 0,  2,  0,  0,  0],
       [ 0, 164,  0,  0,  0],
       [ 0,  1,  0,  0,  1],
       [ 0, 40,  0,  0, 29],
       [ 0,  3,  0,  0, 78]], dtype=int64)
```

In [92]:

```
cm_display = ConfusionMatrixDisplay(confusion_matrix(data_Y_test,
target1_1))
cm_display.plot()
plt.show()
```

No description has been provided for this image

In [93]:

```
# Параметры TP, TN, FP, FN считаются как сумма по всем классам
precision_score(data_Y_test, target1_1, average='micro')
```

Out[93]:

```
0.7610062893081762
```

In [94]:

```
# Параметры TP, TN, FP, FN считаются отдельно для каждого класса
# и берется среднее значение, дисбаланс классов не учитывается.
precision_score(data_Y_test, target1_1, average='macro')
```

```
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\metrics\_classification.py:1509:
```

```
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
```

Out[94]:

```
0.30063492063492064
```

In [95]:

```
classification_report(data_Y_test, target1_1,  
                      target_names=["0", "1", "2", "4", "5"],  
                      output_dict=True)
```

C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\metrics_classification.py:1509:

UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is",  
len(result))
```

C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\metrics_classification.py:1509:

UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is",  
len(result))
```

C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\metrics_classification.py:1509:

UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is",  
len(result))
```

Out[95]:

```
{'0': {'precision': 0.0, 'recall': 0.0, 'f1-score': 0.0, 'support':  
2.0},
```

```
  '1': {'precision': 0.780952380952381,  
        'recall': 1.0,  
        'f1-score': 0.8770053475935828,  
        'support': 164.0},
```

```
  '2': {'precision': 0.0, 'recall': 0.0, 'f1-score': 0.0, 'support':  
2.0},
```

```
  '4': {'precision': 0.0, 'recall': 0.0, 'f1-score': 0.0, 'support':  
69.0},
```

```
  '5': {'precision': 0.7222222222222222,  
        'recall': 0.9629629629629629,  
        'f1-score': 0.8253968253968254,  
        'support': 81.0},
```

```
  'accuracy': 0.7610062893081762,
```

```
  'macro avg': {'precision': 0.30063492063492064,  
                'recall': 0.39259259259259255,  
                'f1-score': 0.34048043459808164,  
                'support': 318.0},
```

```
  'weighted avg': {'precision': 0.5867175801138065,  
                  'recall': 0.7610062893081762,
```

```
'f1-score': 0.6625346536556302,  
'support': 318.0}}
```

In [96]:

```
# Кросс-валидация
```

```
scores = cross_val_score(KNeighborsClassifier(n_neighbors=3), data_x,  
data_y, cv=5)  
scores, np.mean(scores)
```

```
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-  
packages\sklearn\model_selection\_split.py:737: UserWarning: The least  
populated class in y has only 3 members, which is less than  
n_splits=5.  
warnings.warn(
```

Out[96]:

```
(array([0.925      , 0.925      , 0.9375      , 0.93670886, 0.87341772]),  
0.9195253164556962)
```

In [97]:

```
# использование метрики f1  
# https://scikit-learn.org/stable/modules/model\_evaluation.html  
scores = cross_val_score(KNeighborsClassifier(n_neighbors=3),  
data_x, data_y, cv=5,  
scoring='f1_weighted')  
scores, np.mean(scores)
```

```
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-  
packages\sklearn\model_selection\_split.py:737: UserWarning: The least  
populated class in y has only 3 members, which is less than  
n_splits=5.  
warnings.warn(
```

Out[97]:

```
(array([0.91736934, 0.91049719, 0.92523148, 0.93019428, 0.86636917]),  
0.9099322917587654)
```

In [98]:

```
scoring = {'precision': 'precision_weighted',  
          'recall': 'recall_weighted',  
          'f1': 'f1_weighted'}
```

In [99]:

```
scores = cross_validate(KNeighborsClassifier(n_neighbors=2),  
data_x, data_y, scoring=scoring,
```



```

cv=5, return_train_score=True)

scores

C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\model_selection\_split.py:737: UserWarning: The least
populated class in y has only 3 members, which is less than
n_splits=5.
  warnings.warn(
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\metrics\_classification.py:1509:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\metrics\_classification.py:1509:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\metrics\_classification.py:1509:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\metrics\_classification.py:1509:
UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in
labels with no true samples. Use `zero_division` parameter to control
this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\metrics\_classification.py:1509:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))

```

Out[99]:

```
{'fit_time': array([0.00224161, 0.00200868, 0.0010016 , 0.00200391,
0.00199676]),
'score_time': array([0.00909376, 0.00638533, 0.00800085, 0.00843191,
0.00844407]),
'test_precision': array([0.91661491, 0.90184612, 0.91520712,
0.93788637, 0.8672906 ]),
'train_precision': array([0.96285147, 0.96593404, 0.96572028,
0.9693198 , 0.96931968]),
'test_recall': array([0.925      , 0.925      , 0.9375      , 0.93670886,
0.87341772]),
'train_recall': array([0.96226415, 0.96540881, 0.96540881,
0.96865204, 0.96865204]),
'test_f1': array([0.91736934, 0.91049719, 0.92556134, 0.93634402,
0.86513873]),
'train_f1': array([0.96148801, 0.96472415, 0.96508547, 0.96709979,
0.9680539 ])}

```

In [100]:

```
# Grid Search
```

```
n_range = np.array(range(5,35,5))
tuned_parameters = [{'n_neighbors': n_range}]
tuned_parameters
```

Out[100]:

```
[{'n_neighbors': array([ 5, 10, 15, 20, 25, 30])}]
```

In [101]:

```
clf_gs = GridSearchCV(KNeighborsClassifier(), tuned_parameters, cv=5,
scoring='f1_macro')
clf_gs.fit(data_X_train, data_Y_train)
```

```
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\model_selection\_split.py:737: UserWarning: The least
populated class in y has only 1 members, which is less than
n_splits=5.
  warnings.warn(
```

Out[101]:

```
GridSearchCV(cv=5, estimator=KNeighborsClassifier(),
              param_grid=[{'n_neighbors': array([ 5, 10, 15, 20, 25,
30])}],
              scoring='f1_macro')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

GridSearchCV?[Documentation for GridSearchCV](#)iFitted

```
GridSearchCV(cv=5, estimator=KNeighborsClassifier(),
             param_grid=[{'n_neighbors': array([ 5, 10, 15, 20, 25,
30])}],
             scoring='f1_macro')
```

estimator: KNeighborsClassifier

KNeighborsClassifier()

KNeighborsClassifier?[Documentation for KNeighborsClassifier](#)

KNeighborsClassifier()

In [102]:

clf_gs.cv_results_

Out[102]:

```
{'mean_fit_time': array([0.00184278, 0.00126472, 0.00160446,
0.00159974, 0.00152874,
                        0.00148463]),
 'std_fit_time': array([0.00043005, 0.00037938, 0.0004942 ,
0.00049161, 0.00041198,
                        0.00046118]),
 'mean_score_time': array([0.00370574, 0.0039422 , 0.00344634,
0.00314426, 0.00339689,
                        0.00365777]),
 'std_score_time': array([0.00048652, 0.00048739, 0.00077763,
0.00053601, 0.00049155,
                        0.00038674]),
 'param_n_neighbors': masked_array(data=[5, 10, 15, 20, 25, 30],
                                   mask=[False, False, False, False, False, False],
                                   fill_value='?',
                                   dtype=object),
 'params': [{'n_neighbors': 5},
            {'n_neighbors': 10},
            {'n_neighbors': 15},
            {'n_neighbors': 20},
            {'n_neighbors': 25},
            {'n_neighbors': 30}],
 'split0_test_score': array([0.70909091, 0.70909091, 0.70909091,
0.70909091, 0.70909091,
                        0.55614973]),
 'split1_test_score': array([0.83333333, 0.83333333, 0.83333333,
0.83333333, 0.83224401,
                        0.72592593]),
```

```

'split2_test_score': array([0.67222222, 0.67222222, 0.67222222,
0.67222222, 0.58552632,
      0.45      ]),
'split3_test_score': array([0.58552632, 0.73529412, 0.68303571,
0.58552632, 0.45      ,
      0.45      ]),
'split4_test_score': array([0.57222222, 0.57222222, 0.53921569,
0.57222222, 0.45      ,
      0.45      ]),
'mean_test_score': array([0.674479 , 0.70443256, 0.68737957,
0.674479 , 0.60537225,
      0.52641513]),
'std_test_score': array([0.09466447, 0.08497838, 0.0937817 ,
0.09466447, 0.14893172,
      0.10789488]),
'rank_test_score': array([3, 1, 2, 3, 5, 6])}

```

In [103]:

```

# Лучшая модель
clf_gs.best_estimator_

```

Out[103]:

```

KNeighborsClassifier(n_neighbors=10)

```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

KNeighborsClassifier? [Documentation for KNeighborsClassifier](#) [Fitted](#)

```

KNeighborsClassifier(n_neighbors=10)

```

In [104]:

```

# Лучшее значение метрики
clf_gs.best_score_

```

Out[104]:

```

0.7044325609031491

```

In [105]:

```

# Лучшее значение параметров
clf_gs.best_params_

```

Out[105]:

```

{'n_neighbors': 10}

```

In [106]:

```
# Изменение качества на тестовой выборке в зависимости от K-соседей
plt.plot(n_range, clf_gs.cv_results_['mean_test_score'])
```

Out[106]:

```
[<matplotlib.lines.Line2D at 0x20990c85280>]
```

No description has been provided for this image

In [107]:

```
# Randomized Search
```

```
clf_rs = RandomizedSearchCV(KNeighborsClassifier(), tuned_parameters,
cv=5, scoring='f1_macro')
clf_rs.fit(data_X_train, data_Y_train)
```

```
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\model_selection\_search.py:318: UserWarning: The
total space of parameters 6 is smaller than n_iter=10. Running 6
iterations. For exhaustive searches, use GridSearchCV.
```

```
warnings.warn(
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\model_selection\_split.py:737: UserWarning: The least
populated class in y has only 1 members, which is less than
n_splits=5.
warnings.warn(
```

Out[107]:

```
RandomizedSearchCV(cv=5, estimator=KNeighborsClassifier(),
                  param_distributions=[{'n_neighbors': array([ 5, 10,
15, 20, 25, 30])}],
                  scoring='f1_macro')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

RandomizedSearchCV? [Documentation for RandomizedSearchCV](#) [Fitted](#)

```
RandomizedSearchCV(cv=5, estimator=KNeighborsClassifier(),
                  param_distributions=[{'n_neighbors': array([ 5, 10,
15, 20, 25, 30])}],
                  scoring='f1_macro')
```

estimator: KNeighborsClassifier

KNeighborsClassifier()

KNeighborsClassifier?[Documentation for KNeighborsClassifier](#)

```
KNeighborsClassifier()
```

```
In [108]:
```

```
clf_rs.best_score_, clf_rs.best_params_
```

```
Out[108]:
```

```
(0.7044325609031491, {'n_neighbors': 10})
```

```
In [109]:
```

```
plt.plot(n_range, clf_rs.cv_results_['mean_test_score'])
```

```
Out[109]:
```

```
[<matplotlib.lines.Line2D at 0x20990c7a030>]
```

No description has been provided for this image

```
In [110]:
```

```
clf_gs.best_estimator_.fit(data_X_train, data_Y_train)
target2_0 = clf_gs.best_estimator_.predict(data_X_train)
target2_1 = clf_gs.best_estimator_.predict(data_X_test)
```

```
In [111]:
```

```
# Новое качество модели
accuracy_score(data_Y_train, target2_0), accuracy_score(data_Y_test,
target2_1)
```

```
Out[111]:
```

```
(0.85, 0.8459119496855346)
```

```
In [112]:
```

```
# Качество модели до подбора гиперпараметров
accuracy_score(data_Y_train, target2_0), accuracy_score(data_Y_test,
target2_1)
```

```
Out[112]:
```

```
(0.85, 0.8459119496855346)
```

```
In [113]:
```

```
# Оценка качества
accuracy_score(data_Y_train, target2_0), accuracy_score(data_Y_test,
target2_1)
```

```
Out[113]:
```

```
(0.85, 0.8459119496855346)
```

```
In [114]:
```

```
confusion_matrix(data_Y_test, target2_1)
```

```
Out[114]:
```

```
array([[ 0,  2,  0,  0,  0],
       [ 0, 152,  0, 12,  0],
       [ 0,  1,  0,  0,  1],
       [ 0, 11,  0, 44, 14],
       [ 0,  0,  0,  8, 73]], dtype=int64)
```

```
In [115]:
```

```
cm_display = ConfusionMatrixDisplay(confusion_matrix(data_Y_test,
target2_1))
cm_display.plot()
plt.show()
```

No description has been provided for this image

```
In [116]:
```

```
# Параметры TP, TN, FP, FN считаются как сумма по всем классам
precision_score(data_Y_test, target2_1, average='micro')
```

```
Out[116]:
```

```
0.8459119496855346
```

```
In [117]:
```

```
# Параметры TP, TN, FP, FN считаются отдельно для каждого класса
# и берется среднее значение, дисбаланс классов не учитывается.
precision_score(data_Y_test, target2_1, average='macro')
```

```
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\metrics\_classification.py:1509:
```

```
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
```

```
Out[117]:
```

```
0.4865416210295728
```

```
In [118]:
```

```
classification_report(data_Y_test, target2_1,
                      target_names=["0", "1", "2", "3", "5"],
                      output_dict=True)
```

```
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\metrics\_classification.py:1509:
```

```
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
```

```
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\metrics\_classification.py:1509:
```

```
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
```

```
C:\Users\Yan note\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\metrics\_classification.py:1509:
```

```
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0
in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
```

```
Out[118]:
```

```
{'0': {'precision': 0.0, 'recall': 0.0, 'f1-score': 0.0, 'support':
2.0},
```

```
 '1': {'precision': 0.9156626506024096,
       'recall': 0.926829268292683,
       'f1-score': 0.9212121212121213,
       'support': 164.0},
```

```
 '2': {'precision': 0.0, 'recall': 0.0, 'f1-score': 0.0, 'support':
2.0},
```

```
 '3': {'precision': 0.6875,
       'recall': 0.6376811594202898,
       'f1-score': 0.6616541353383458,
       'support': 69.0},
```

```
 '5': {'precision': 0.8295454545454546,
       'recall': 0.9012345679012346,
       'f1-score': 0.863905325443787,
       'support': 81.0},
```

```
 'accuracy': 0.8459119496855346,
```

```
 'macro avg': {'precision': 0.4865416210295728,
               'recall': 0.4931489991228415,
               'f1-score': 0.48935431639885085,
               'support': 318.0},
```

```
 'weighted avg': {'precision': 0.8327023789842044,
                  'recall': 0.8459119496855346,
```



```
'f1-score': 0.838708347729813,  
'support': 318.0}}
```

In [118]: